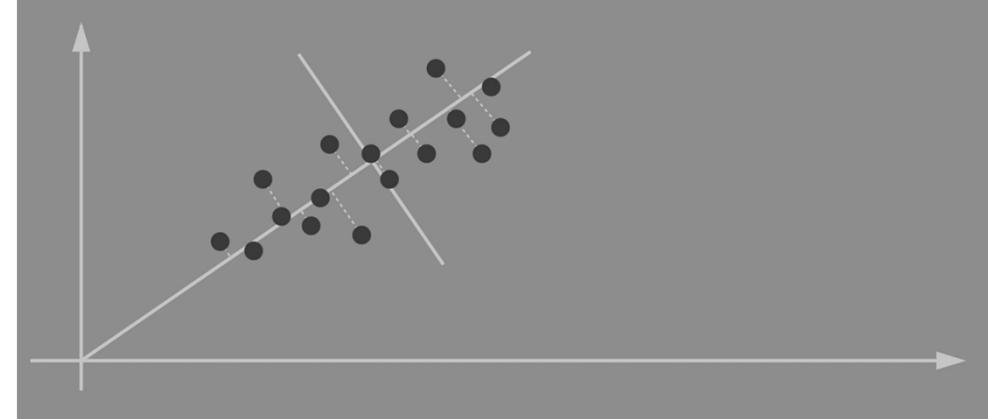




https://i0.wp.com/thedatascientist.com/wp-content/uploads/2018/05/recommender_systems.png



[PCA-06-scaled.jpg \(2560×1051\) \(perfectial.com\)](#)

Lecture 8: Scalable Matrix Factorisation for Collaborative Filtering in RecSys & Scalable PCA for Dimensionality Reduction

COM6012: Scalable ML by Haiping Lu

Week 8 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

Week 8 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

Many Decisions to Make



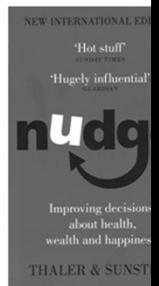
Recommendations Everywhere

Frequently bought



- This item: Thinking, F...
- Atomic Habits: The lif...
- How to Win Friends a...

Customers who



Nudge: Improving Decisions About Health, Wealth and Happiness
Richard H. Thaler

job

- jobs at amazon
- job in sheffield
- jobs near me
- jobs sheffield
- jobs ecclesall
- jobs ac uk
- jobs at amazon
- Jobcentre Plus
- jobs indeed uk

Online events for you



Scholars Webinar on Drug Discovery

Wed, Mar 24 - Thu,

Alireza Khanteymo...

View

People you may know



Twin Karma...

Research Software
9 mutual connections

Connect

ChatGPT 4o ▾



Can you recommend a book for me to read this summer?

Can you recommend a good book to read this year?

Can you recommend a good book to read this month?

Can you recommend interesting documentaries to watch?

Can you recommend



Search



Deep research



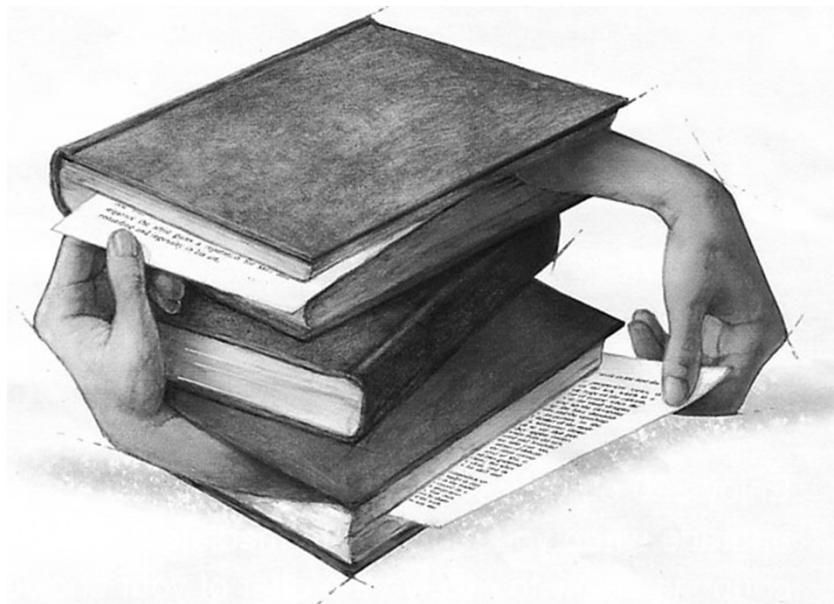
Recommender Systems (RecSys)

- Predict relevant items for a user, in a given context
- Predict to what extent these items are relevant
- A ranking task (searching as well)
- Implicit, targeted, **intelligent** advertisement
- Effective, popular marketing

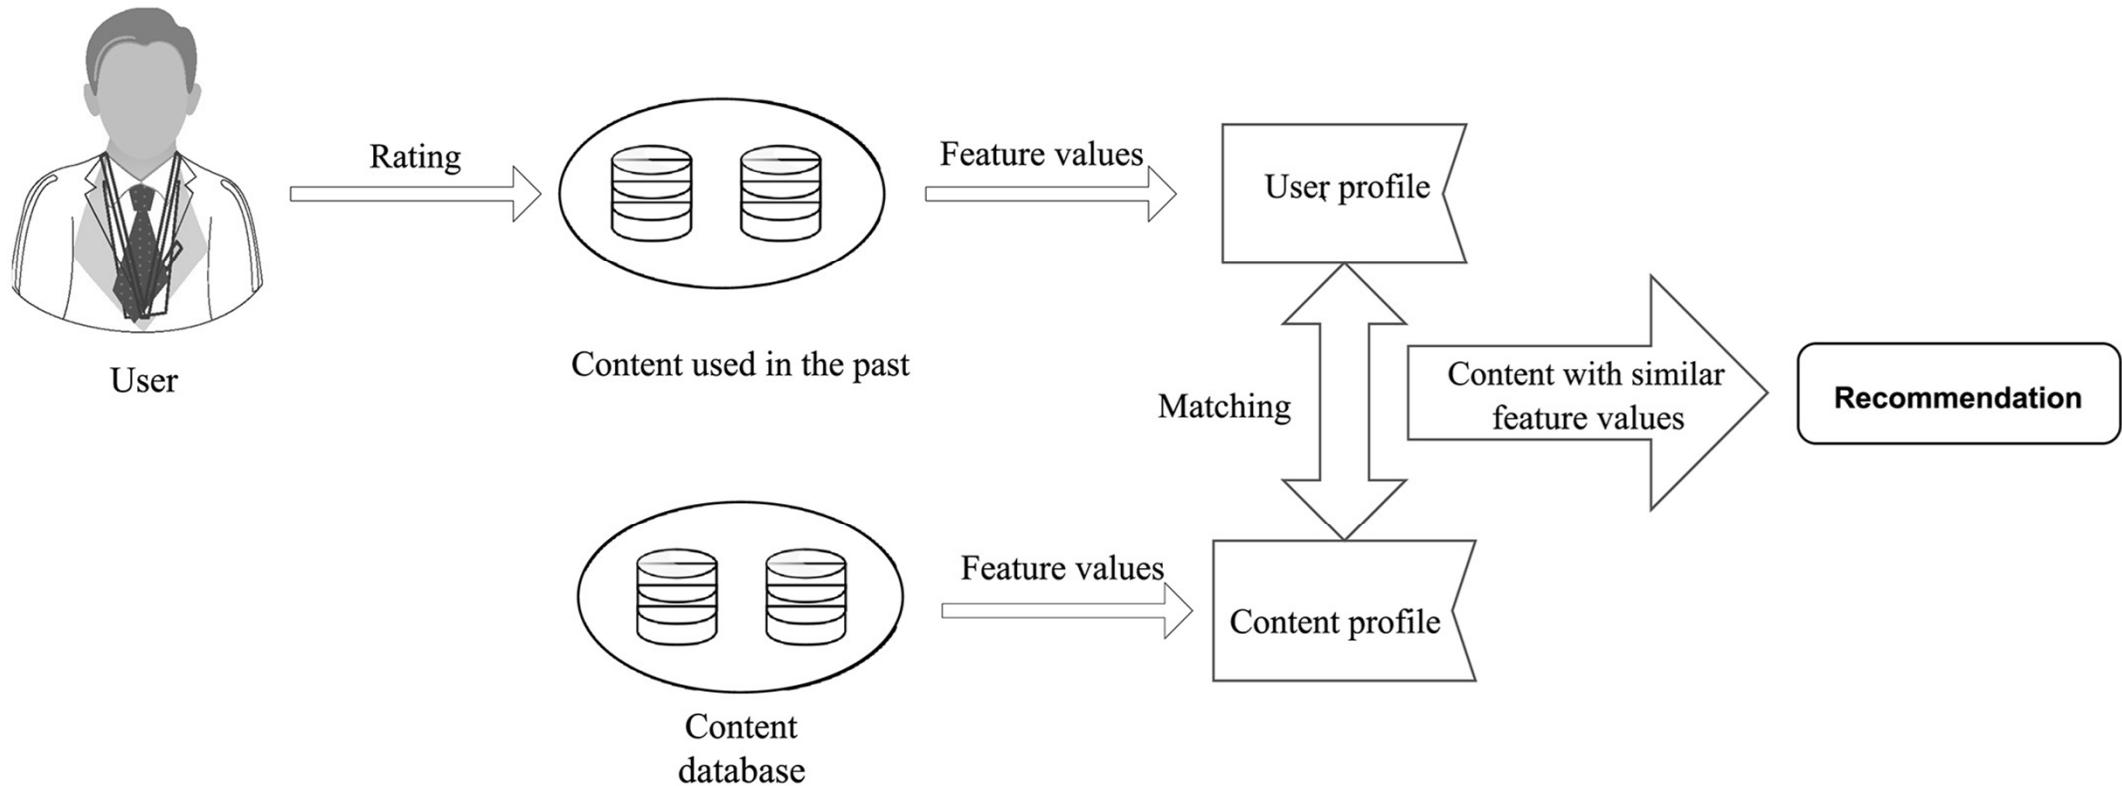


Two Classes of RecSys

- Content-based recommender systems
- Collaborative filtering recommender systems

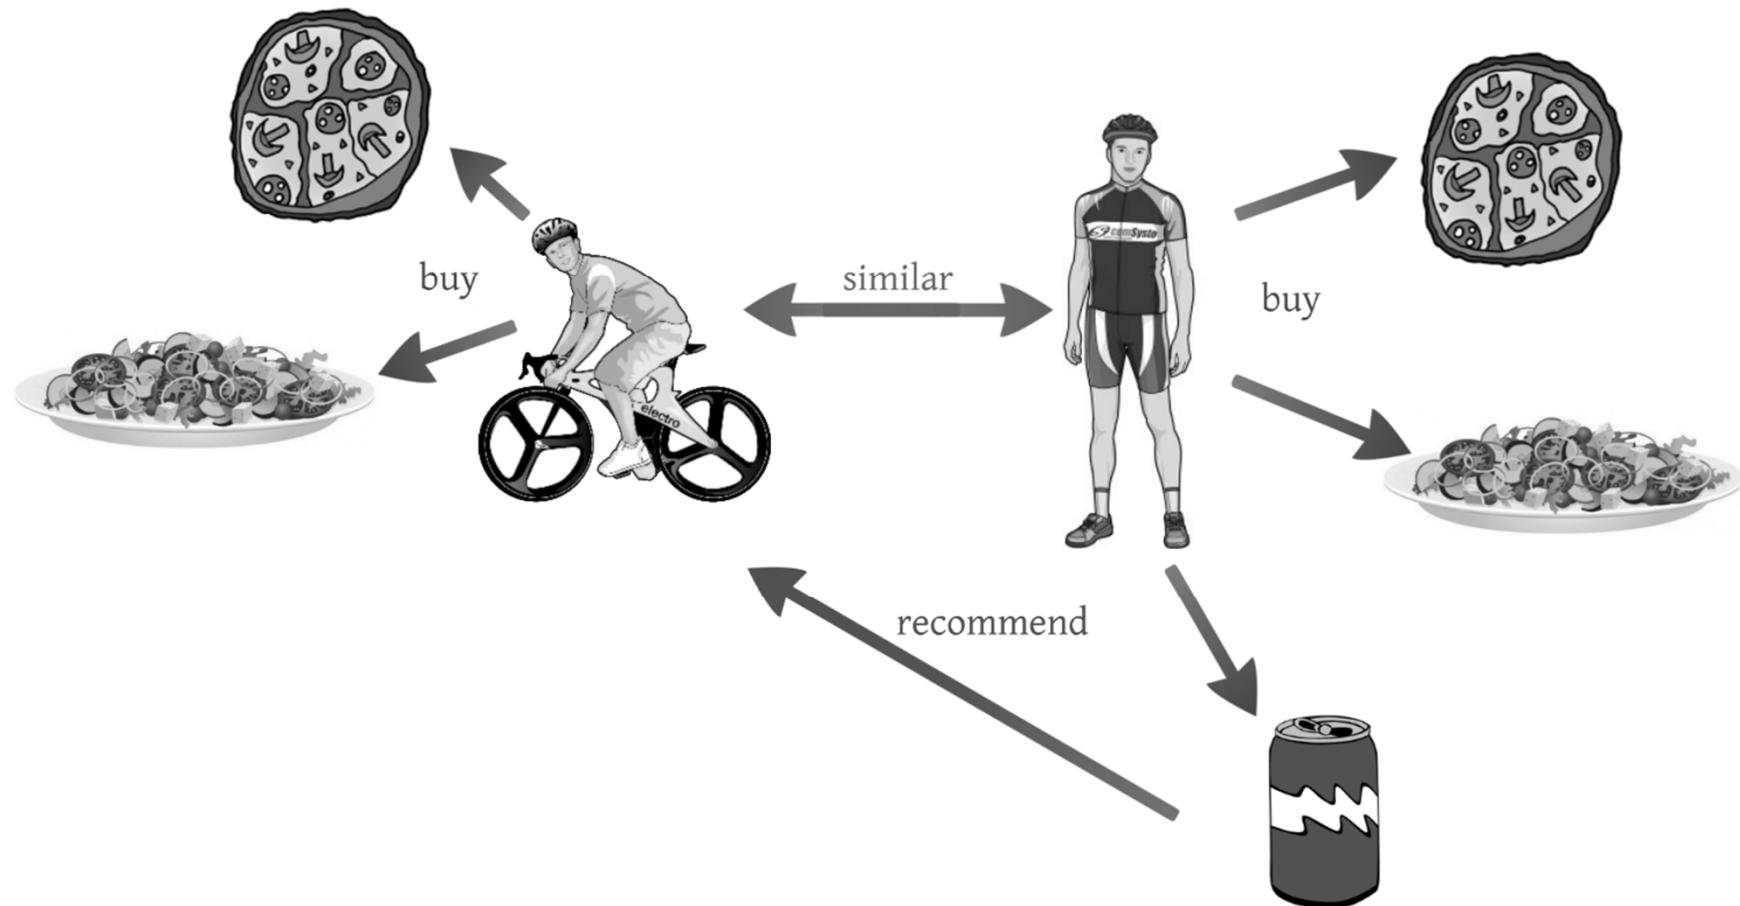


Content-based RecSys



[dac4058-fig-0010-m.jpg \(2128x789\) \(wiley.com\)](#)

Collaborative Filtering RecSys



https://miro.medium.com/max/2656/1*6_NIX6CJYhtxzRM-t6ywkQ.png

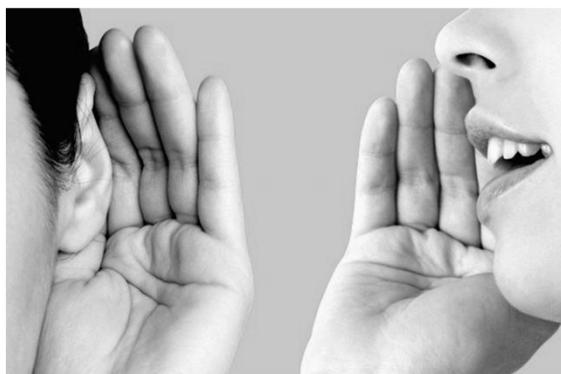
Week 8 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

What is Collaborative Filtering?

- Information filtering based on past records
 - Electronic word of mouth marketing
 - Turn visitors into customers (e-Salesman)
- Components
 - Users (customers): who provide ratings
 - Items (products): to be rated
 - Ratings (interest): core data



John	5	1	3	5
Tom	?	?	?	2
Alice	4	?	3	?

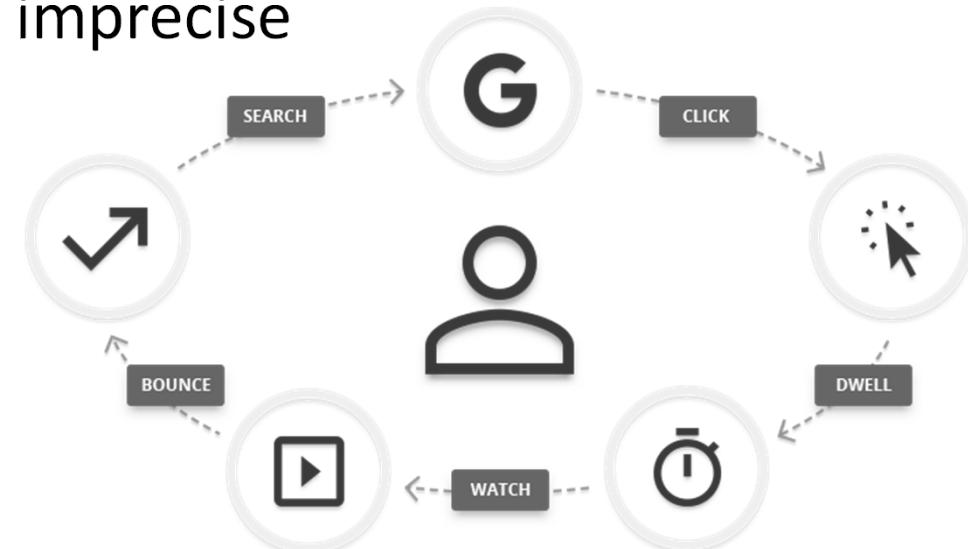
Collaborative Filtering (CF)

- Objective: predict how well a user will like an unrated item, given past ratings for a community of users
- How does CF work?
 - Input: many users' ratings for many items
 - Model: similar users \leftarrow ratings strongly correlate
 - Recommend items rated highly by similar users



Explicit vs Implicit Ratings

- Explicit (direct): users indicate levels of interest
 - Most accurate descriptions of a user's preference
 - Challenging in collecting data
- Implicit (indirect): observing user behavior
 - Can be collected with little or no cost to user
 - Ratings inference may be imprecise



Rating Scales

- Scalar ratings
 - Numerical scales
 - 1-5, 1-7, etc.
- Binary ratings
 - Agree/Disagree, Good/Bad, etc.
- Unary ratings
 - Presence/absence of an event, e.g., purchase/browsing history, search patterns, mouse movements
 - No info about the opposite $\neq 0$



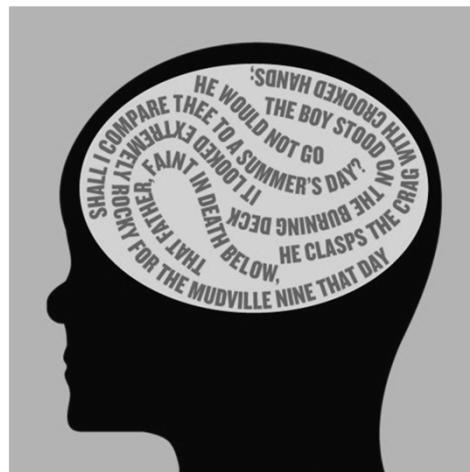
CF Preferences

- Many users, many items, many ratings
- Users rate multiple items
- Other users with similar needs/tastes
- Item evaluation requires personal taste
- Taste persists



CF Methods

- Memory-based: predict using past ratings directly
 - Weighted ratings given by other similar users
 - User-based & item-based (non-ML)
- Model-based: model users based on past ratings
 - Predict ratings using the learned model



[iforget-465.jpg \(465x465\) \(newyorker.com\)](#)



[neural-header.jpg \(756x503\) \(utsouthwestern.edu\)](#)

Prediction Accuracy

- Mean absolute error (MAE)

$$MAE = \frac{\sum_{i,j} |p_{i,j} - r_{i,j}|}{n}$$

- Normalized MAE

$$NMAE = \frac{MAE}{r_{max} - r_{min}}$$

- Root mean squared error (RMSE)

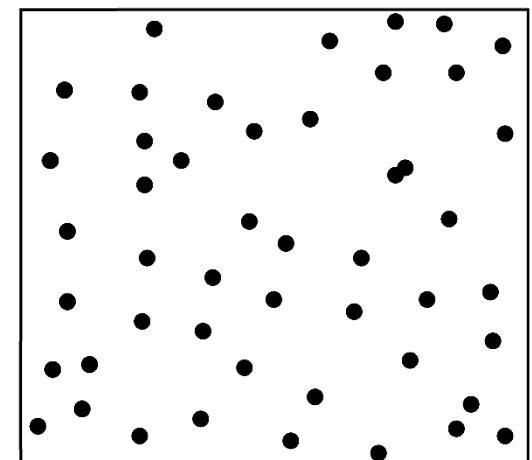
$$RMSE = \sqrt{\frac{1}{n} \sum_{i,j} (p_{i,j} - r_{i,j})^2}$$

Challenges

- Cold Start
 - New user
 - Rate some initial items
 - Non-personalized rec.
 - Describe tastes
 - Demographic info
 - New item
 - Randomly selecting items
 - Content analysis, metadata (non-CF)
- Sparsity: sparse user-item matrix
- Scalability: millions of users and items



[isbil+2.jpg \(490x303\) \(squarespace-cdn.com\)](#)



Week 8 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

Matrix Factorisation (MF) for CF

- Characterise items/users by vectors of factors learned from the rating matrix user x item
- High correlation between item and user factors → good recommendation
- Flexibility: incorporate implicit feedback, temporal effects, and confidence levels



	Argo	Seven	Inception	Interstellar
John	5	1	3	5
Tom	?	?	?	2
Alice	4	?	3	?

Basic MF Model

- Map users & items to a joint latent factor space of dimensionality k
 - Item $i \rightarrow$ vector q_i : the extent to which the item possesses those k factors
 - User u : vector p_u : the extent of interest the user has on those k factors
- User-item interactions: the user's overall interest in the item's characteristics
 - Inner product $q_i^T p_u$: predicted user u 's rating of item i

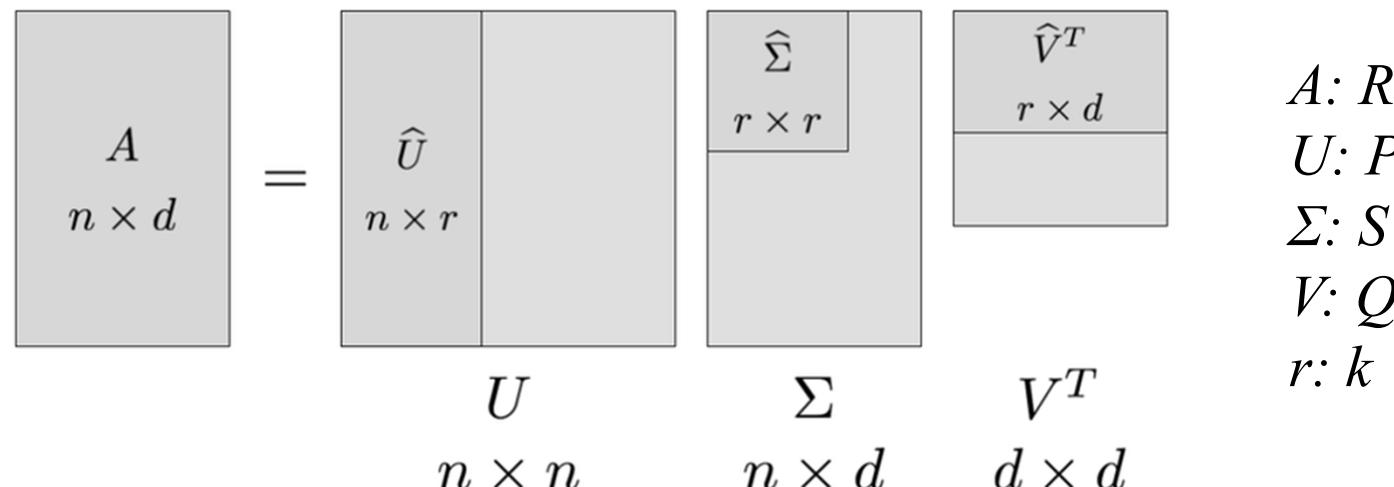
$$\hat{r}_{u,i} = q_i^T p_u$$

How to Learn the MF Model

- To learn: item factors $\{q_i\}$ and user factors $\{p_u\}$
- Factorisation assuming full rating matrix
 - Factorise rating matrix R using SVD to obtain P, S, Q

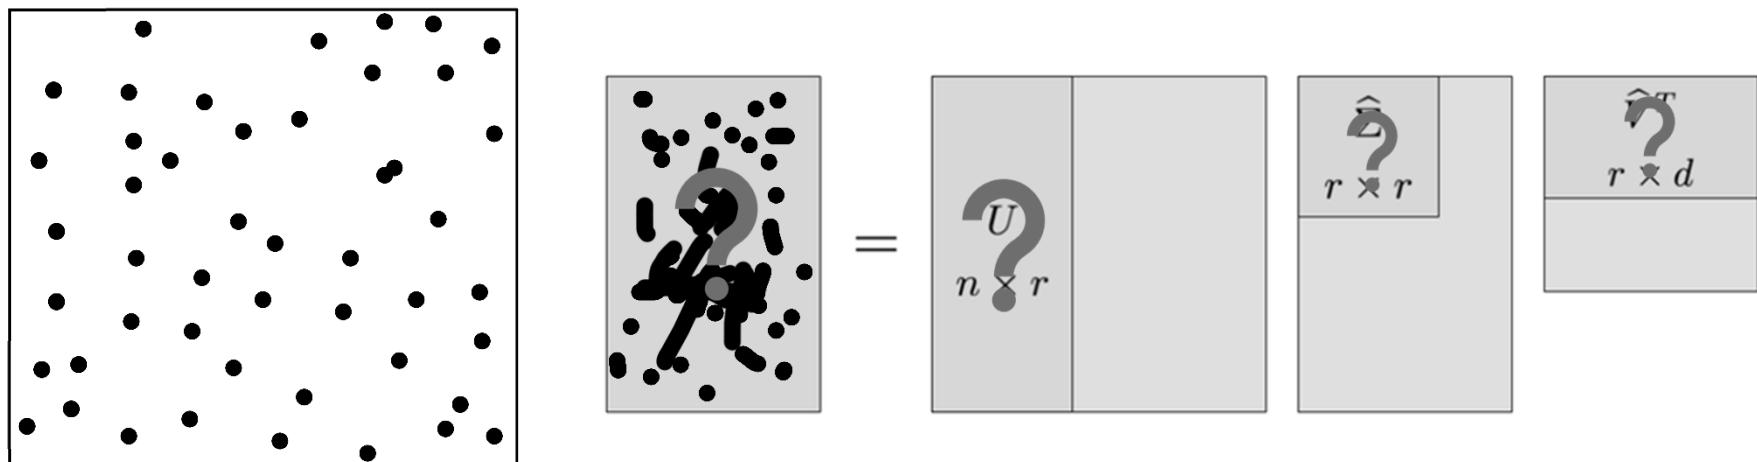
$$R = PSQ^T$$

- Reduce the matrix S to dimension k , i.e. S_k
- $P \rightarrow P_k$ and $Q \rightarrow Q_k$: $P_k S_k \rightarrow \hat{P}$, and $S_k Q_k^T \rightarrow \hat{Q}^T$
- u th row of $\hat{P} \rightarrow p_u$, i th column of $\hat{Q}^T \rightarrow q_i$



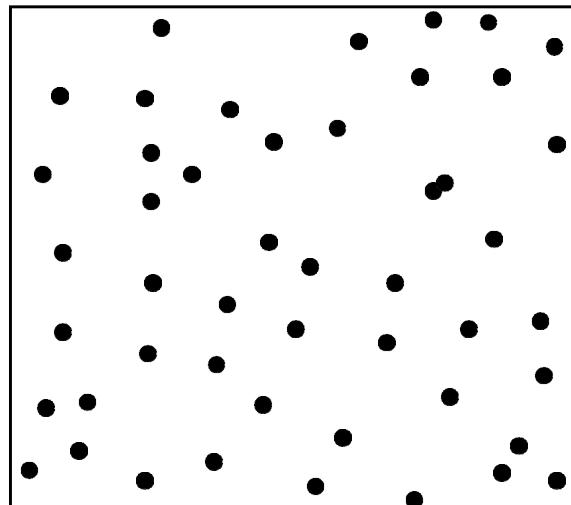
Challenges in MF for CF

- High portion of missing values caused by sparseness in the user-item rating matrix
- Conventional SVD is undefined when knowledge about the matrix is incomplete



How to Fill Missing Values

- Imputation: fill in missing ratings using the average ratings for user and item
- Problems
 - Expensive: significantly increases the amount of data
 - Inaccurate imputation might distort the data



MF with Missing Values

- Modelling directly the observed ratings only
 - Avoid overfitting through a regularised model
 - Minimize the regularised squared error on the set of known ratings to learn the factor vectors p_u and q_i

$$\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2)$$

- κ : the training set of the (u,i) pairs with known ratings
- λ : the regularisation parameter

Alternating Least Squares for MF-CF

$$\min_{q^*, p^*} \sum_{(u,i) \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\| q_i \|_2^2 + \| p_u \|_2^2)$$

- Both p_u and q_i are unknown (non-convex function)
- Fix one of them → quadratic with optimal solution
- Alternating Least Squares (ALS): alternate between fixing q_i s and fixing p_u s
 - Fix $P(p_u)$ s as \hat{P} to recompute q_i s by solving a least-squares problem $\| R - PQ^T \|_F$ (Frobenius norm)

$$R = \hat{P}Q^T \Rightarrow Q^T = (\hat{P}^T \hat{P})^{-1} \hat{P}^T R$$

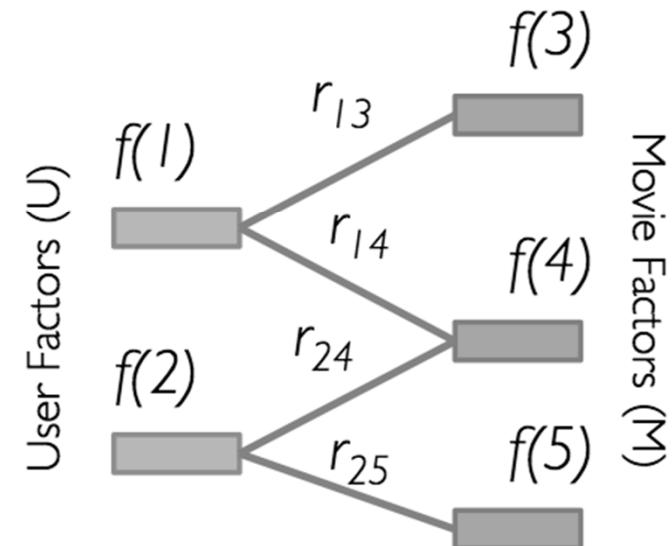
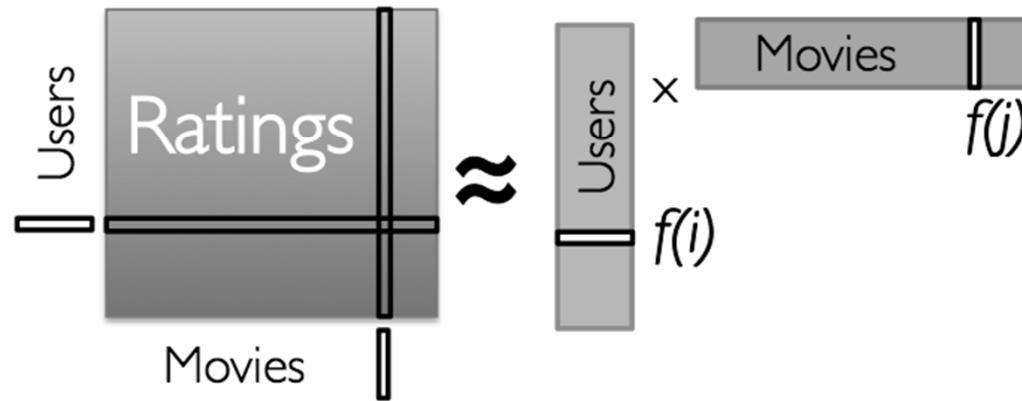
- Fix Q as \hat{Q} , we have

$$P = R\hat{Q}(\hat{Q}^T \hat{Q})^{-1}$$

- Random initialisation to start this iteration

MF for Movie Recommendation

Low-Rank Matrix Factorization:



Iterate:

$$f[i] = \arg \min_{w \in \mathbb{R}^d} \sum_{j \in \text{Nbrs}(i)} (r_{ij} - w^T f[j])^2 + \lambda ||w||_2^2$$

Figure [fro](#)_{spark-training/matrix_factorization.png} at master · databricks/spark-training (github.com)

Week 8 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

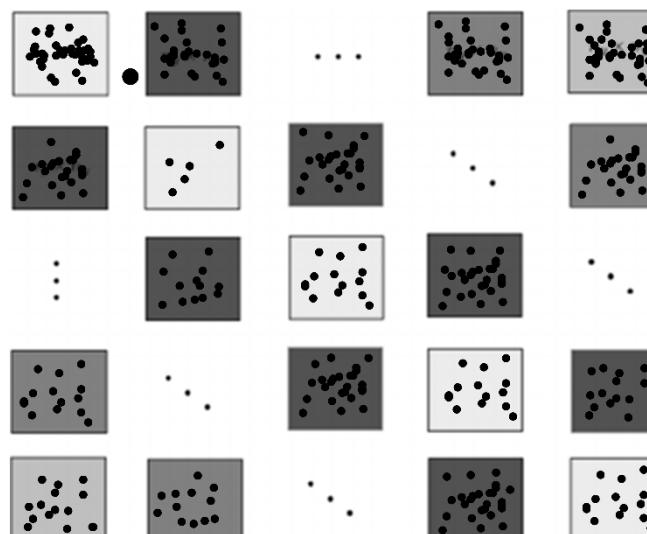
- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

Key in Scalable ML

- Computation and storage should be linear (in n, d)
→ Low-cost computation (time + space)
- Perform parallel and in-memory computation
→ Many working + reduce disk I/O
- Minimise network communication → Reduce overhead in parallelisation, not the more the better

Blocked Implementation of ALS

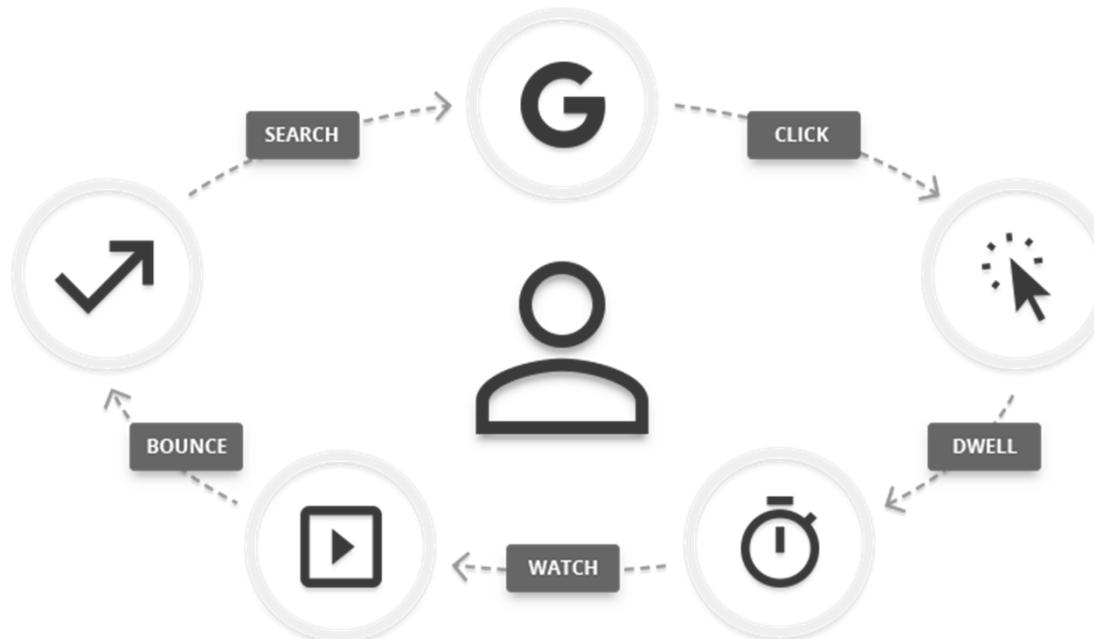
- Group users and items into blocks
 - Reduce communication: only send one copy of each user vector to each item block on each iteration, and only for the item blocks that need that user's feature vector
 - Pre-compute info: out-links of each user (which blocks of items it will contribute to); in-links for each item (which of the feature vectors it receives from each user block it will depend on)



[Color-online-A-symmetric-block-Toeplitz-matrix-Each-block-is-also-a-symmetric-Toeplitz.png \(488x369\) \(researchgate.net\)](#)

Implicit Feedback Modelling

- Implicit feedback: views, clicks, purchases, likes, shares
 - Rating $r =$ strength in observations of user actions (#clicks, viewing duration) → confidence level in observed user preference
 - Construct a preference matrix P : e.g. 1 if $r > 0$ and 0 if $r = 0$
 - Factorisation of P → latent factors to predict the preference of a user for an item (details in an [ICDM08 paper](#))



The ALS API in Spark

- numUserBlocks/numItemBlocks: the number of blocks the users/items will be partitioned into to parallelize computation (defaults to 10)
- rank: the number of latent factors in the model (defaults to 10)
- regParam: the regularization parameter in ALS (defaults to 0.1)
- implicitPrefs: whether to use the explicit feedback ALS variant or one adapted for implicit feedback data (defaults to false: explicit ratings)
- alpha: the baseline confidence of implicit feedback (defaults to 1.0)
- nonnegative: whether to use nonnegative constraints (defaults to false)
- coldStartStrategy: “drop” → drop any rows in the DataFrame of predictions that contain NaN values (defaults to “nan”: assign NaN to a user and/or item factor is not present in the model)
- blockSize: the size of the user/product blocks in the blocked implementation of ALS to reduce communication

CF in Spark ML

- Scala code (1800+ lines)
- Documentation: Collaborative Filtering in Spark
- DataBricks movie recommendations tutorial
- DataBricks: founded by the creators of Apache Spark
 - Their latest packages at their GitHub page
 - Databricks community edition: free with limited features (14-day free trial of the full platform).



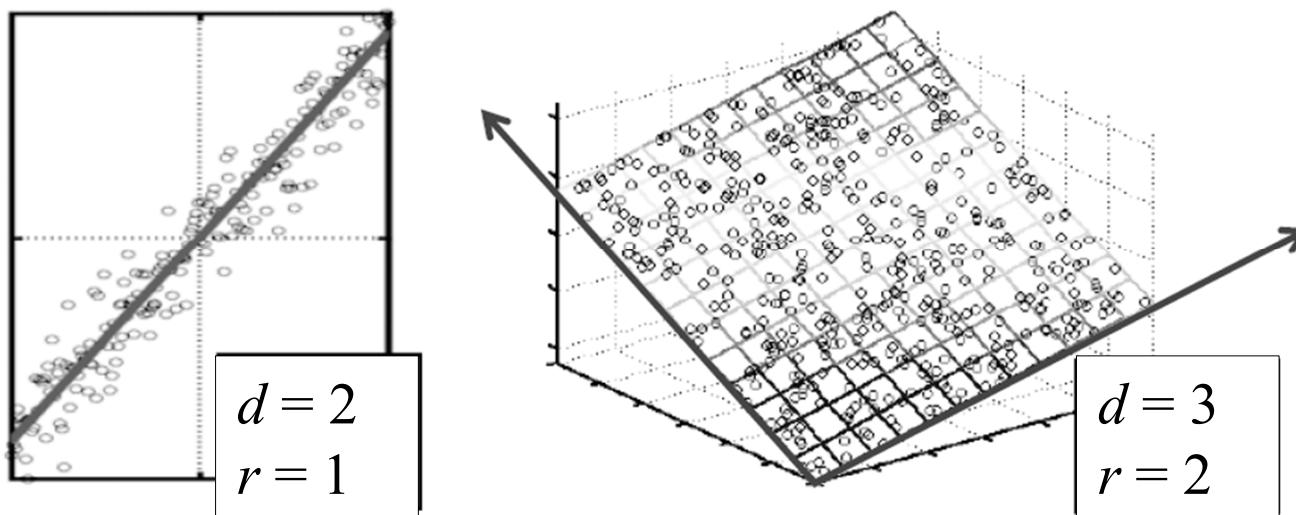
Week 8 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

Motivation of Dimensionality Reduction

- Raw data: complex and high-dimensional
- Assumption: they lie on a low-dimensional subspace
 - Axes of this subspace → representation of the data
 - Simpler, more compact, showing interesting patterns



Utilities of Dimensionality Reduction

- Discover hidden correlations/topics
- Remove redundant/noisy features
- Interpretation and visualisation
- Easier storage and processing of the data



[owners-icebergs-blog-image-300x300.jpg \(resettogrow.com\)](#)



[1*KvKlx9OnIxdoTfNxWKAY_g.jpeg \(480x320\) \(medium.com\)](#)



[Interpreting and Translation Blog: Image \(wordpress.com\)](#)

PCA → Variance Maximisation

- Input: n d -dimensional data points
- PCA algorithm
 - $X_0 \leftarrow n \times d$ data matrix, data point \rightarrow row vector x_i
 - X : subtract mean x from each row vector x_i in X_0
 - $\Sigma \leftarrow X^T X$: Gramian/scatter matrix for X
 - Find eigenvectors and eigenvalues of Σ
 - $U (d \times r) \leftarrow$ the top r eigenvectors (PCs)
- PCA features for y : $U^T y$ (dimension: $d \rightarrow r$)
 - Zero correlations, ordered by variance

Scalability Problem of PCA

- Input dimensionality → scatter matrix
 - Images: $100 \times 100 \rightarrow 10^4$; $1000 \times 1000 \rightarrow 10^6$
 - Scatter matrix Σ is of size d^2
 - $d = 10^4 \rightarrow \Sigma$ size 10^8
 - $d = 10^6 \rightarrow \Sigma$ size $= 10^{12}$
- Alternative: Singular Value Decomposition (SVD)
 - Efficient algorithms available
 - Often need just top r eigenvectors

Week 8 Contents / Objectives

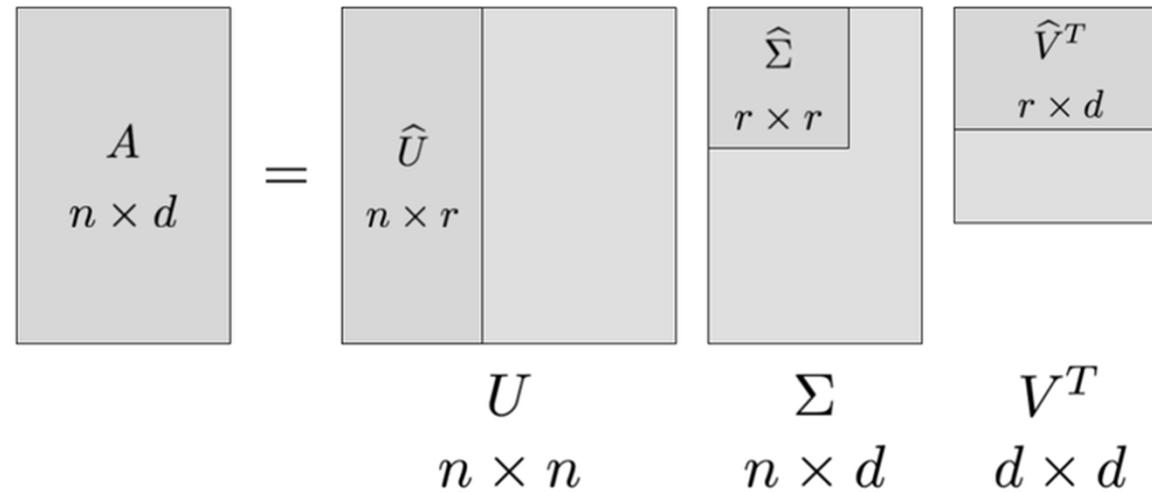
- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

Singular Value Decomposition (SVD)

$$A_{[n \times d]} = U_{[n \times r]} \Sigma_{[r \times r]} (V_{[d \times r]})^T$$

- A: $n \times d$ matrix
- r: the rank of the matrix
- U: $n \times r$ matrix, column orthonormal, $U^T U = I$
- Σ : $r \times r$ diagonal matrix, strength of each factor
- V: $d \times r$ matrix, column orthonormal, $V^T V = I$



SVD \leftrightarrow Eigen-decomposition

- SVD gives
 - $X = U \Sigma V^T$
- Eigen-decomposition gives
 - $B = X^T X = W \Lambda W^T$
- U, V, W : orthonormal $\rightarrow U^T U = I, V^T V = I, W^T W = I$
- Σ, Λ : diagonal
- Relationship:
 - $XX^T = U \Sigma V^T (U \Sigma V^T)^T = U \Sigma V^T (V \Sigma^T U^T) = U \Sigma^2 U^T$
 - $X^T X = V \Sigma^T U^T (U \Sigma V^T) = V \Sigma \Sigma^T V^T = V \Sigma^2 V^T$
 - $B = X^T X = W \Lambda W^T = V \Sigma^2 V^T$

PCA via SVD

- $X_0 \leftarrow n \times d$ data matrix, data point \rightarrow row vector x_i
- X : subtract mean x from each row vector x_i in X_0
- $U \Sigma V^T \leftarrow$ SVD of X
- The top r right singular vectors V of $X \rightarrow$ the PCs
- The singular values in $\Sigma =$ the square roots of the eigenvalues of $X^T X$

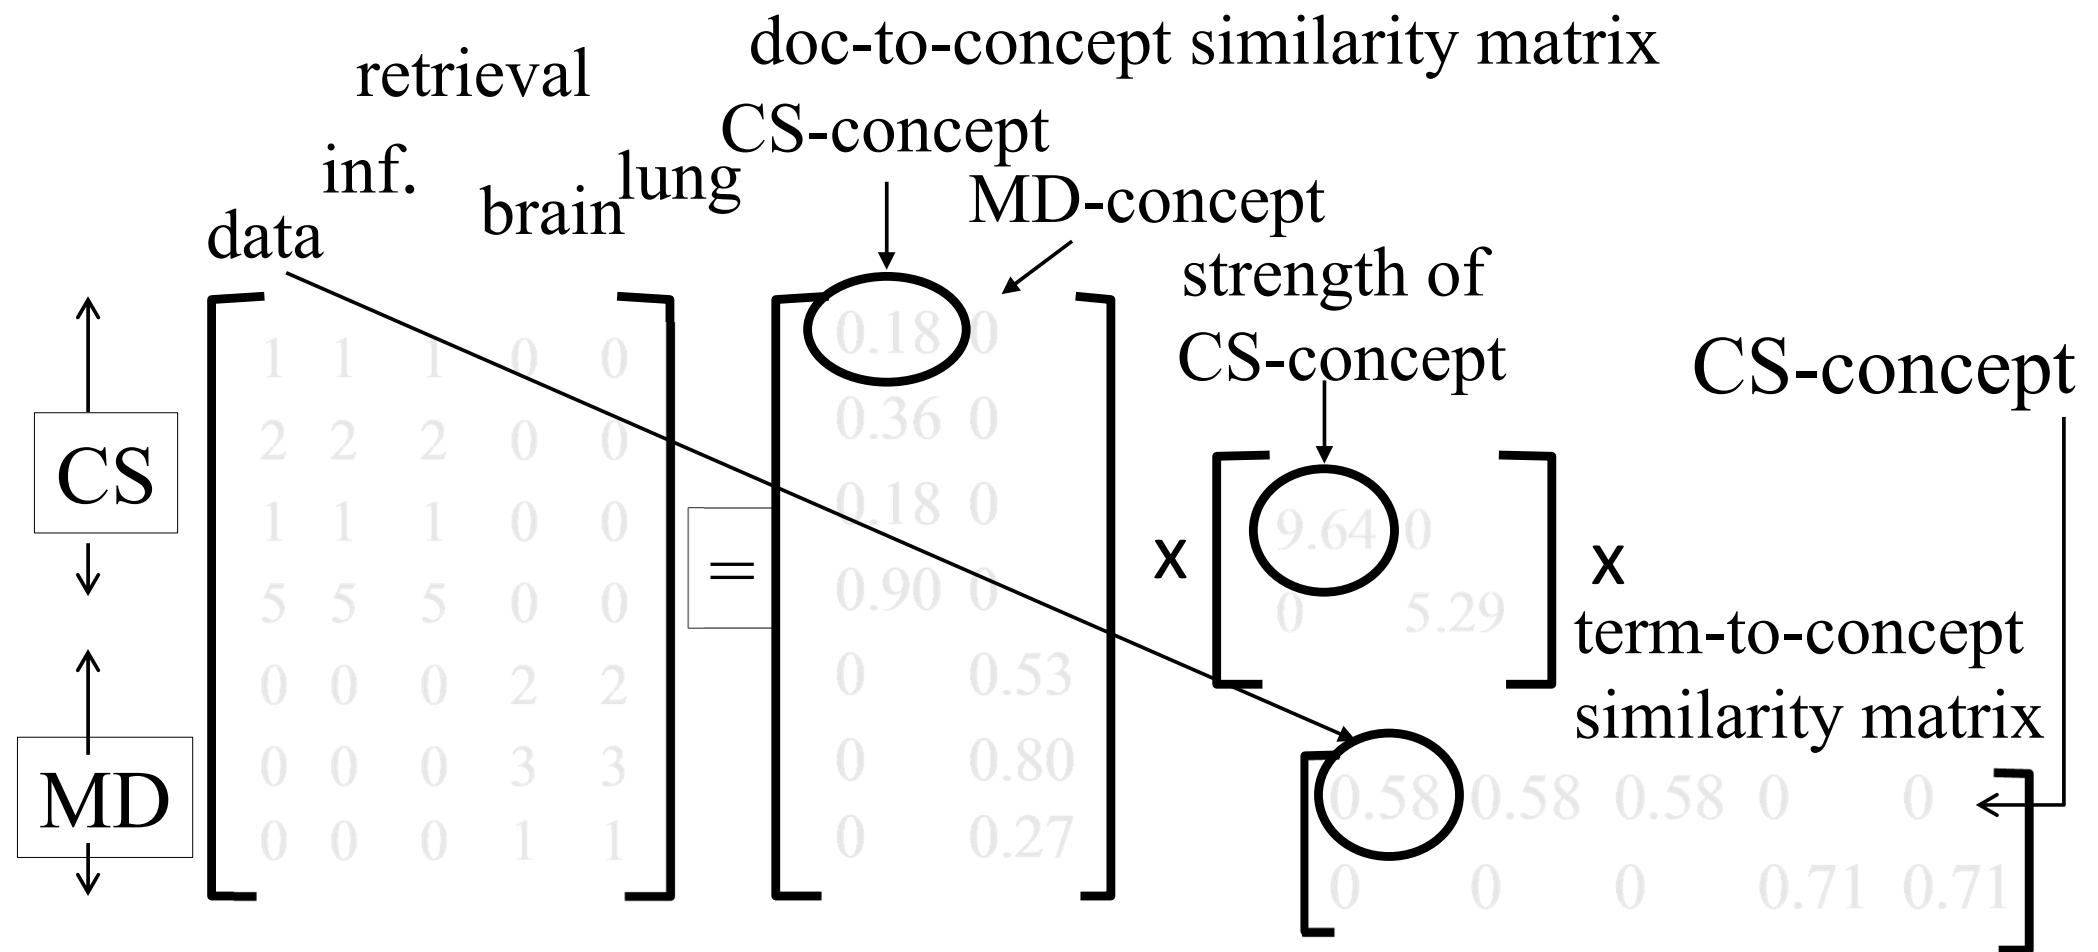
Example on a Document x Term

Term Document	data	information	retrieval	brain	lung
CS-TR1	1	1	1	0	0
CS-TR2	2	2	2	0	0
CS-TR3	1	1	1	0	0
CS-TR4	5	5	5	0	0
MED-TR1	0	0	0	2	2
MED-TR2	0	0	0	3	3
MED-TR3	0	0	0	1	1

- $d = 5$ but $r=2 \rightarrow$ two bases $[1 1 1 0 0]$ & $[0 0 0 1 1]$
- U: document-to-concept similarity matrix
- V: term-to-concept similarity matrix
- Σ : its diagonal elements \rightarrow strength of each concept

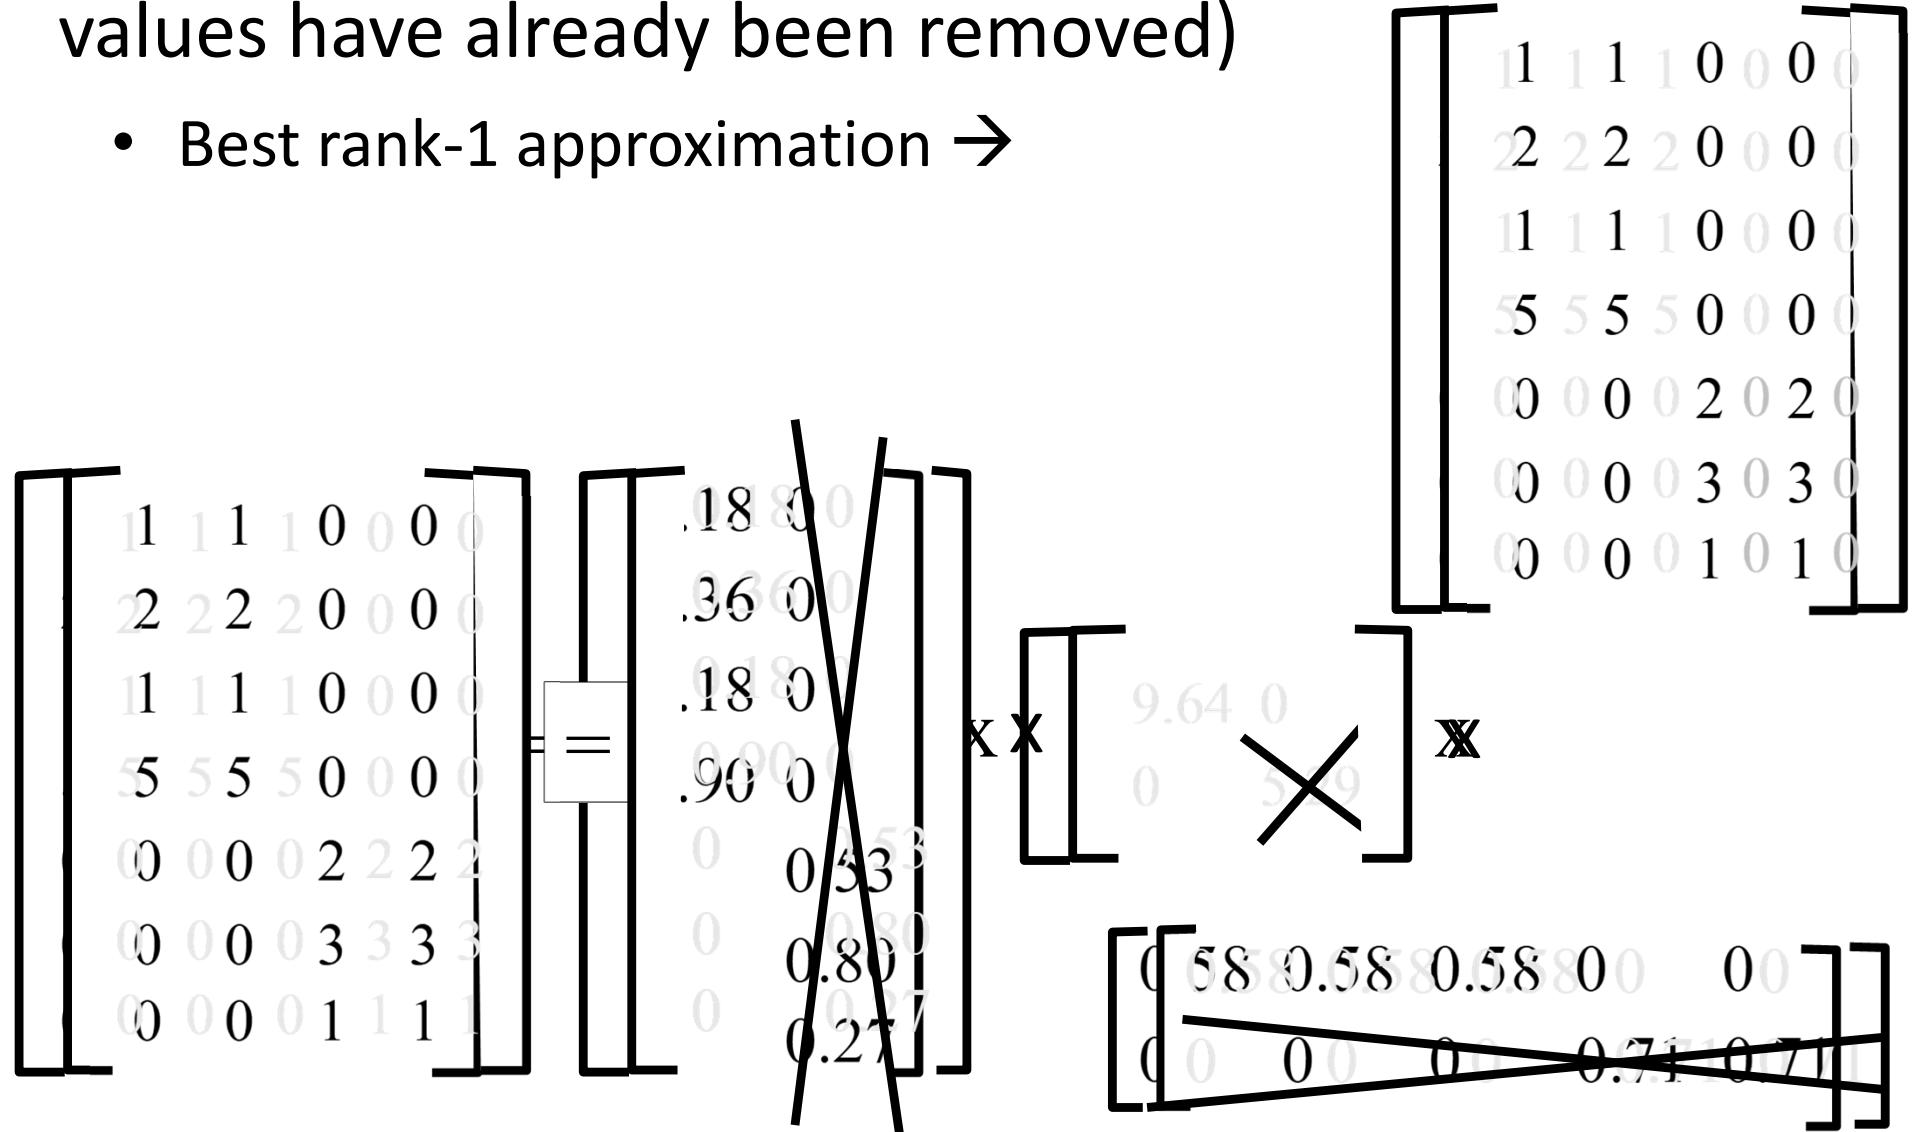
Interpretation

Term Document	data	information	retrieval	brain	lung
CS-TR1	1	1	1	0	0
CS-TR2	2	2	2	0	0
CS-TR3	1	1	1	0	0
CS-TR4	5	5	5	0	0
MED-TR1	0	0	0	2	2
MED-TR2	0	0	0	3	3
MED-TR3	0	0	0	1	1



SVD – Dimensionality Reduction

- To reduce the dimensionality further (3 zero singular values have already been removed)
 - Best rank-1 approximation →



Week 8 Contents / Objectives

- Recommender Systems
- Collaborative Filtering
- Matrix Factorisation for Collaborative Filtering
- Scalable Collaborative Filtering in Spark

- Scalability Problem of PCA
- PCA via SVD
- Scalable PCA in Spark

Three PCA APIs in Spark MLlib

- DataFrame-based API – PCA (source code, Scala doc)
 - `pyspark.ml.feature.PCA(k=None, inputCol=None, outputCol=None)`
- RDD-based API – RowMatrix (source code, Scala doc)
 - `computePrincipalComponents(k)`
 - **Scalable:** `computeSVD(k, computeU=False, rCond=1e-09)`

```
@Since("1.6.0")
def computePrincipalComponentsAndExplainedVariance(k: Int): (Matrix, Vector) = {
    val n = numCols().toInt
    require(k > 0 && k <= n, s"k = $k out of range (0, n = $n]")
    if (n > 65535) {
        val svd = computeSVD(k)
        val s = svd.s.toArray.map(eigValue => eigValue * eigValue / (n - 1))
        val eigenSum = s.sum
        val explainedVariance = s.map(_ / eigenSum)
    }
}
```

SVD in Spark MLlib (RDD)

- $U: m \times k$; $\Sigma : k \times k$; $V: n \times k$
- Assumption: n (dimensionality) < m (# samples)
- Different methods based on computational cost
 - If n is small ($n < 100$) or k is large compared with n ($k > n/2$), compute $A^T A$ first and then compute its top eigenvalues and eigenvectors locally on the driver
 - Otherwise, compute $(A^T A)v$ in a distributive way and send it to ARPACK to compute $(A^T A)$'s top eigenvalues/eigenvectors on the driver node

Selection of SVD Computation

```
if (n < 100 || (k > n / 2 && n <= 15000)) {
    // If n is small or k is large compared with n, we better compute the Gramian matrix first
    // and then compute its eigenvalues locally, instead of making multiple passes.
    if (k < n / 3) {
        SVMode.LocalARPACK
    } else {
        SVMode.LocalLAPACK
    }
} else {
    // If k is small compared with n, we use ARPACK with distributed multiplication.
    SVMode.DistARPACK
}

case "local-svd" => SVMode.LocalLAPACK
case "local-eigs" => SVMode.LocalARPACK
case "dist-eigs" => SVMode.DistARPACK
case _ => throw new IllegalArgumentException("Do not support mode $mode.")
```

Acknowledgement & References

- Acknowledgement
 - Some slides are adapted from the MMDS book slides.
- References
 - Yehuda Koren, Robert Bell, and Chris Volinsky. "Matrix factorization techniques for recommender systems." *Computer* 8 (2009): 30-37 (Yahoo & AT&T)
 - Yifan Hu, Yehuda Koren, and Chris Volinsky. "Collaborative filtering for implicit feedback datasets." *Eighth IEEE International Conference on Data Mining*, 2008
 - Charu C. Aggarwal, Recommender Systems: The Textbook, April 2016
 - Chapter 11 of the MMDS book